

# Project Architecture: Black-Box Stiff/Non-Stiff ODE Solver (FRANKENSTEIN)

We will build an automatic Julia-based solver for ordinary differential equations (ODEs) that handles both stiff and non-stiff problems transparently. The solver should accept a user's problem (ODE function, initial conditions, time span, optional events) and internally decide the best integration strategy. In particular, it will **detect stiffness** on-the-fly and **switch solvers** as needed, handle **events/discontinuities** via callbacks, and expose a simple "black-box" API with minimal user intervention. Our design will leverage Julia's high-performance ODE ecosystem while adding logic to glue components together.

## Stiff vs Non-Stiff ODEs

An ODE is informally "*stiff*" if explicit integration methods (like Runge–Kutta) require prohibitively small steps to remain stable. There is no single definition of stiffness, but common indicators include multiple time scales (fast decays plus slow dynamics), a large stiffness ratio (ratio of largest to smallest eigenvalues), or outright failure of explicit solvers <sup>1</sup>. Stiff problems typically demand implicit or semi-implicit methods (which use Jacobians) for efficiency, while non-stiff problems can use faster explicit integrators. Our solver must handle both regimes **without requiring the user to classify the problem**.

## Solver Methods and Automatic Switching

We will incorporate both **explicit** and **implicit** integrators, and automatically switch between them as needed. For non-stiff phases, high-order explicit Runge–Kutta methods (e.g. Tsitouras 5(4) – *Tsit5*) are very efficient <sup>2</sup>. For stiff regimes, we'll use stiff solvers like Rosenbrock methods (e.g. *Rosenbrock23*) or implicit BDF-type methods (e.g. *TRBDF2*, *Rodas5P*) that handle stability robustly <sup>3</sup>. In practice, an *auto-switch* algorithm will start with an explicit solver and monitor indicators of stiffness; if stiffness is detected, it will restart or continue with an implicit solver. For example, DifferentialEquations.jl's "AutoTsit5(Rosenbrock23())" integrator uses an explicit Tsit5 stepper until stiffness criteria trigger a switch to Rosenbrock23 <sup>4</sup> <sup>5</sup>. Similarly, algorithms like LSODA (from ODEPACK) always start non-stiff (Adams) and switch to stiff (BDF) when needed <sup>6</sup>. Our design will follow these models: treat stiffness detection and switching as internal to the solver, hiding it from the user.

- **Built-in Auto-Switch Algorithms:** We will implement a two-stage integrator that pairs a non-stiff and a stiff method (e.g. Tsit5 and Rosenbrock23) under an `AutoSwitch` scheme. This mirrors SciML's approach, which provides configurable parameters (such as counts of consecutive stiffness indicators and tolerance thresholds) to control switching <sup>7</sup>. For instance, we may count how many recent steps violate stability limits (`maxstiffstep`) and switch to the stiff solver after a threshold <sup>7</sup>. Conversely, if the stiff solver succeeds with large steps for many steps, we may switch back. Tunable factors (like step-size factor `dtfac` when switching) ensure smooth transitions <sup>7</sup>. Using these heuristics, the solver remains robust across regimes.
- **Recommended Methods:** Based on literature, we plan to include or interface to the following:

- Non-stiff: high-order explicit Runge–Kutta (e.g. *Tsit5*, *BS5*, *Vern7*, *Vern9*) for efficiency and accuracy on smooth non-stiff ODEs <sup>2</sup>.
- Stiff: singly-diagonally-implicit RK (SDIRK) or Rosenbrock methods (e.g. *Rosenbrock23*, *KenCarp4*, *Rodas5P*) which have good stability on stiff problems <sup>8</sup>.
- Specialized: multi-step stiff solvers like BDF (via *CVODE\_BDF* in Sundials.jl) or classical Radau IIA integrators if needed for very large or highly stiff systems.
- Auto-choose: a mechanism to combine the above, e.g. `AutoVern7(Rodas5())` for high-accuracy cases <sup>5</sup>.

The solver's default will use one of the proven auto-switchers (e.g. `AutoTsit5(Rosenbrock23())`), which “works if you know nothing about the equation” <sup>5</sup>. This meets the requirement of not intruding on the user's choice.

## Stiffness Detection

To avoid user intervention, the solver will attempt to **detect stiffness dynamically**. Possible strategies include: - **Step behavior monitoring**: If an explicit integrator fails to progress (e.g. step gets rejected repeatedly, step size collapses) this signals stiffness. We would then restart with an implicit method. - **Stability region check**: Some integrators (like explicit methods with embedded Jacobian estimates) can estimate stability limits. We can compare the current step size or derivative ratio against a threshold to flag stiffness. - **Jacobians and eigenvalues**: Optionally, compute (or approximate) the Jacobian of `f` and inspect its eigenvalue spread. A large stiffness ratio (max/min eigenvalue) is a direct indicator. This is costly for large systems, but doable for moderate-sized problems.

Any chosen detector will feed into the switching logic (like the `nonstifftol` and `stifftol` thresholds in SciML's `AutoSwitch` <sup>7</sup>). In practice, we will implement an internal routine that periodically evaluates stiffness (e.g. every few steps) and triggers a solver change when needed. All such complexity remains hidden; the user simply solves their ODE, and the solver under the hood dispatches the appropriate method.

## Event and Callback Handling

The solver must support **events/discontinuities** robustly. We will adopt a callback-based approach: the user can supply *continuous events* (functions of `(t,y)` whose zero-crossing triggers a state change) and *discrete events* (predicate conditions on `(t,y)` checked at each step). Internally: - **Continuous events**: We will use root-finding on the event functions to locate the exact time `t_event` when an event condition goes to zero. When `condition(t,y)=0` is detected (sign change), we backtrack the last step and solve a nonlinear equation (e.g. via bisection or secant) to find `t_event` to user tolerance. Then we apply the event's effect (update state or parameters) before proceeding <sup>9</sup>. - **Discrete events**: After each successful step, we evaluate any discrete predicates. If one is true, we immediately apply its effect (e.g. jump in state) and continue. No root find is needed, but we must ensure stability (the solver step may be rolled back if a discrete event changes continuity).

This callback mechanism is standard in advanced ODE solvers <sup>9</sup>. We will design an `EventManager` in our architecture to manage these: it keeps lists of event conditions and effects, monitors them, and interrupts the integrator when triggered. After each event, we ensure consistency (e.g. if parameters change, we update derivatives or Jacobians). All of this will be automatic once events are registered, so the user experience remains simple.

## User Interface and Automation

The front-end will expose a single function (e.g. `solve_ode(problem; options...)`) where `problem` includes the ODE function `f(t,y)`, initial conditions, time span, tolerances, and any events. The solver should: 1. **Choose or accept an algorithm implicitly.** By default it will use an auto-switch integrator (like `AutoTsit5(Rosenbrock23())`), unless the user provides hints (e.g. an `alg_hint=:stiff` keyword). We will rely on internal logic to set defaults intelligently <sup>10</sup> <sup>5</sup>. 2. **Require minimal user input.** The user should not need to specify whether the problem is stiff. Optional hints (like tolerances or maximum step) can be provided, but we assume none by default. 3. **Allow Jacobian specification.** If the user can optionally supply a Jacobian function, we'll use it to speed up implicit solves. Otherwise we will compute Jacobians via automatic differentiation (e.g. `ForwardDiff.jl`) or finite differences. For example, SciML's Rosenbrock solvers use autodiff by default <sup>11</sup>, which we can emulate. We will expose an option like `jac=true/false` or accept a Jacobian function in the problem definition. 4. **Return a solution object.** Including time points, solution trajectory (with interpolation between steps), and any event records. This object should be user-friendly (like SciML's `ODESolution`) but we will focus on correctness and speed.

The key is **invisibility of complexity**: users simply call our solver, and it “just works” on stiff or non-stiff ODEs, handling events, without needing to pick an algorithm. Behind the scenes, our code will instantiate the appropriate integrators and callbacks and manage their interaction.

## Performance and Optimizations

Speed is a priority, so we will optimize common bottlenecks: - **In-place operations:** We will use Julia's in-place function calls (e.g. modify a preallocated array `du`) to avoid allocations each step. This follows best practice for high-performance ODE solvers. - **Multi-threading:** For large systems (hundreds of ODEs), we can enable multi-threading in our integrators. Some methods (e.g. explicit Runge-Kutta) allow internal parallel evaluation of `f`. We will exploit Julia's threading when beneficial. (`DifferentialEquations.jl` suggests switching to multi-threading above ~150 ODEs <sup>12</sup>.) - **Vectorization and BLAS:** Use array operations and BLAS for linear algebra in implicit solves. If a Jacobian is dense, we solve linear systems using optimized libraries. - **Sparse/Matrix-free:** If problems are large and sparse, we can attempt to use sparse Jacobians or matrix-free Krylov solvers (like GMRES). We will design our linear solver interface to allow plugging in sparse or iterative methods if the user's problem demands it. - **Tuning step sizes:** The auto-switch algorithm itself may shrink or grow the step size when switching (parameter `dtfac` in `AutoSwitch`) <sup>7</sup>. We'll set sensible defaults (e.g. doubling or halving step size upon switch) but allow user control if needed. - **Use of existing libraries:** Rather than coding all solvers from scratch, we can call efficient Julia implementations. For example, using `OrdinaryDiffEq.jl`'s solvers internally (invoking them as algorithms within our wrapper) will give us battle-tested performance. If absolutely needed for very large stiff systems, we can interface `Sundials.jl` or `LSODA.jl`. Indeed, SciML documentation recommends using `LSODA` or `CVODE_BDF` for >1000 equations <sup>4</sup>. We will design the architecture so that replacing or adding solver backends (Julia or external) is straightforward.

## Development Plan Summary

In summary, our architecture will include:

- **Problem Preprocessor:** Validates user input, sets up ODE system structure, and chooses initial solver strategy (e.g. default auto-switcher).

- **Integrator Modules:** Julia implementations of explicit and implicit ODE solvers (or wrappers to them). E.g. Tsit5(), Rosenbrock23(), TRBDF2(), etc. Each module exposes a stepper interface.
- **Stiffness Detector & Switcher:** Wraps two integrators. Monitors solution at each step, applies logic (counters and tolerances) to switch between solvers as in an AutoSwitch algorithm <sup>7</sup>.
- **Event Manager/Callbacks:** Handles continuous and discrete events. Ties into the integrators to stop, find zero-crossings, and apply state changes <sup>9</sup>.
- **Solver Interface:** A high-level function (e.g. `solve_ode`) that glues everything together, applies default options, and returns a solution object.

Each component will be thoroughly tested (with standard ODE test suites, including stiff/unstiff examples and event problems). We will document the design in the final README, describing how each piece works and how to use the solver.

**Sources:** This design follows best practices in modern ODE solvers <sup>4</sup> <sup>5</sup>. It leverages automatic method selection (as in SciML's auto-switchers) and callback-based event handling <sup>9</sup>, aiming for a fast, user-friendly Julia package. By building on these proven strategies and making them transparent to the user, FRANKENSTEIN will be a flexible yet powerful “black-box” ODE solver for both stiff and non-stiff systems.

---

<sup>1</sup> What is considered a stiff ordinary differential equation (ODE) in MATLAB 8.0 (R2012b)? - MATLAB Answers - MATLAB Central

<https://www.mathworks.com/matlabcentral/answers/94354-what-is-considered-a-stiff-ordinary-differential-equation-ode-in-matlab-8-0-r2012b>

<sup>2</sup> <sup>3</sup> <sup>4</sup> <sup>7</sup> <sup>8</sup> <sup>10</sup> <sup>11</sup> <sup>12</sup> ODE Solvers · DifferentialEquations.jl

[https://docs.sciml.ai/DiffEqDocs/stable/solvers/ode\\_solve/](https://docs.sciml.ai/DiffEqDocs/stable/solvers/ode_solve/)

<sup>5</sup> Getting Started with Differential Equations in Julia · DifferentialEquations.jl

[https://docs.sciml.ai/DiffEqDocs/stable/getting\\_started/](https://docs.sciml.ai/DiffEqDocs/stable/getting_started/)

<sup>6</sup> Isoda: Solver for Ordinary Differential Equations (ODE), Switching... in deSolve: Solvers for Initial Value Problems of Differential Equations ('ODE', 'DAE', 'DDE')

<https://rdrr.io/cran/deSolve/man/Isoda.html>

<sup>9</sup> Event Handling and Callback Functions · ModelingToolkit.jl

<https://docs.sciml.ai/ModelingToolkit/stable/basics/Events/>